



# GAN - Theory and Applications

---

Michele de Simoni

Paolo Galeone

Emanuele Ghelfi

Federico di Mattia

March 13, 2019



# Generative Adversarial Networks

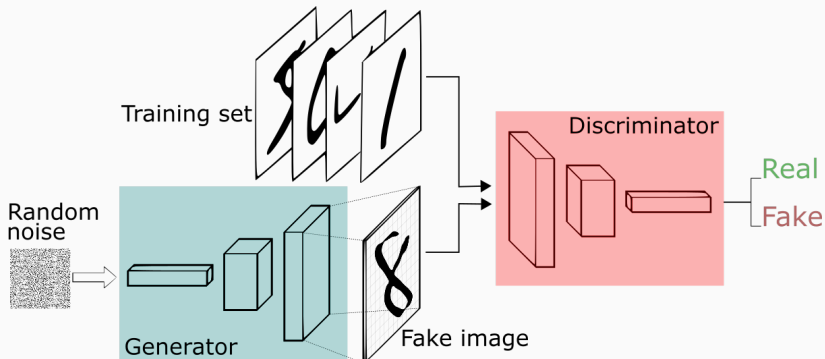
“Adversarial Training (also called GAN for Generative Adversarial Networks) is the most interesting idea in the last 10 years of ML.”

— Yann LeCun

# Generative Adversarial Networks

Two components, the **generator** and the **discriminator**:

- The **generator**  $G$ , aim is to capture the data distribution.
- The **discriminator**  $D$ , estimates the probability that a sample came from the training data rather than from  $G$ .



**Figure 1:** Credits: Reference

# Generative Adversarial Networks

Generator and Discriminator compete against each other, playing the following **zero sum min-max game** with value function  $V_{GAN}(D, G)$ :

$$\min_G \max_D V_{GAN}(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \quad (1)$$

# Generative Adversarial Networks

Generator and Discriminator compete against each other, playing the following **zero sum min-max game** with value function  $V_{GAN}(D, G)$ :

$$\min_G \max_D V_{GAN}(D, G) = \underbrace{\mathbb{E}_{x \sim p_{data}(x)} [\log D(x)]}_{\text{real samples}} + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \quad (1)$$

# Generative Adversarial Networks

Generator and Discriminator compete against each other, playing the following **zero sum min-max game** with value function  $V_{GAN}(D, G)$ :

$$\min_G \max_D V_{GAN}(D, G) = \underbrace{\mathbb{E}_{x \sim p_{data}(x)} [\log D(x)]}_{\text{real samples}} + \underbrace{\mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]}_{\text{generated samples}} \quad (1)$$

Intuitive explanation:

- **Discriminator** needs to:
  - Correctly classify real data:

$$\max_D \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)]. \quad (2)$$

- Correctly classify wrong data:

$$\max_D \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]. \quad (3)$$

Intuitive explanation:

- **Generator** needs to **fool** the discriminator:
  - Generate samples similar to the real one:

$$\min_G \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]. \quad (4)$$



# GANs - Generator

Intuitive explanation:

- **Generator** needs to **fool** the discriminator:
  - Generate samples similar to the real one:


$$\min_G \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]. \quad (4)$$

- Saturates easily (1).
- Change loss for generator:

$$\max_G \mathbb{E}_{z \sim p_z(z)} [\log(D(G(z)))]. \quad (5)$$


# GANs - Models definition

- Both D and G can be parametrized functions (Neural Networks).
- Different architectures for different data types.
  - Tuple of numbers?



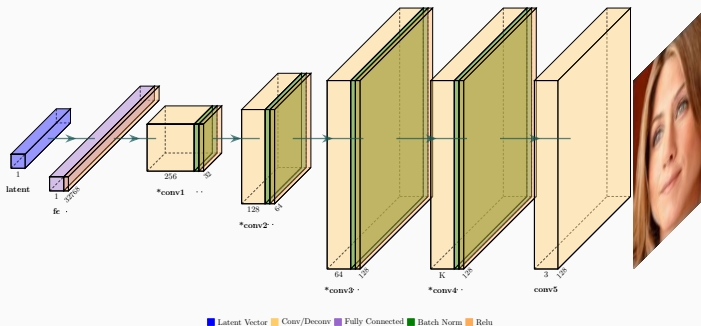
## GANs - Models definition

- Both D and G can be parametrized functions (Neural Networks).
- Different architectures for different data types.
  - Text or sequences?



# GANs - Models definition

- Both D and G can be parametrized functions (Neural Networks).
- Different architectures for different data types.
  - Images?



# GAN Training

# GANs - Training

- Discriminator and generator are **competing** against each other.
- How to train?
- **Alternating** execution of training steps.
- Use **minibatch stochastic gradient descent/ascent**.



How to **train** the **discriminator**?

Repeat from 1 to **k**:

1. Sample minibatch of  $m$  noise samples  $z^{(1)}, \dots, z^{(m)}$  from noise prior  $p_g(z)$
2. Sample minibatch of  $m$  examples  $x^{(1)}, \dots, x^{(m)}$  from data generating distribution  $p_{data}(x)$
3. Train the **discriminator** by stochastic gradient **ascent**:

$$\Delta_{\theta_d} \frac{1}{m} \sum_{i=1}^m \log D(x^{(i)}) + \log(1 - D(G(z^{(i)})))$$

# GANs - Training - Discriminator

How to **train** the **discriminator**?

Repeat from 1 to **k**:

1. Sample minibatch of  $m$  noise samples  $z^{(1)}, \dots, z^{(m)}$  from noise prior  $p_g(z)$
2. Sample minibatch of  $m$  examples  $x^{(1)}, \dots, x^{(m)}$  from data generating distribution  $p_{data}(x)$
3. Train the **discriminator** by stochastic gradient **ascent**:

$$\Delta_{\theta_d} \frac{1}{m} \sum_{i=1}^m \underbrace{\log D(x^{(i)}) + \log(1 - D(G(z^{(i)})))}_{\text{Discriminator loss}}$$



# GANs - Training - Discriminator

How to **train** the **discriminator**?

Repeat from 1 to **k**:

1. Sample minibatch of  $m$  noise samples  $z^{(1)}, \dots, z^{(m)}$  from noise prior  $p_g(z)$
2. Sample minibatch of  $m$  examples  $x^{(1)}, \dots, x^{(m)}$  from data generating distribution  $p_{data}(x)$
3. Train the **discriminator** by stochastic gradient **ascent**:

$$\Delta_{\theta_d} \underbrace{\frac{1}{m} \sum_{i=1}^m \underbrace{\log D(x^{(i)}) + \log(1 - D(G(z^{(i)})))}_{\text{Discriminator loss}}}_{\text{Loss estimation using m samples}}$$

How to **train** the **generator**?

The update is executed **only once** and only after the turn of the discriminator is completed:

1. Sample minibatch of  $m$  noise samples  $z^{(1)}, \dots, z^{(m)}$  from noise prior  $p_g(z)$
2. Train the **generator** by stochastic gradient **ascent**:

$$\Delta_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(D(G(z^{(i)})))$$

How to **train** the **generator**?

The update is executed **only once** and only after the turn of the discriminator is completed:

1. Sample minibatch of  $m$  noise samples  $z^{(1)}, \dots, z^{(m)}$  from noise prior  $p_g(z)$
2. Train the **generator** by stochastic gradient **ascent**:

$$\Delta_{\theta_g} \frac{1}{m} \sum_{i=1}^m \underbrace{\log(D(G(z^{(i)})))}_{\text{Generator loss}}$$

How to **train** the **generator**?

The update is executed **only once** and only after the turn of the discriminator is completed:

1. Sample minibatch of  $m$  noise samples  $z^{(1)}, \dots, z^{(m)}$  from noise prior  $p_g(z)$
2. Train the **generator** by stochastic gradient **ascent**:

$$\Delta_{\theta_g} \underbrace{\frac{1}{m} \sum_{i=1}^m \underbrace{\log(D(G(z^{(i)})))}_{\text{Generator loss}}}_{\text{Loss estimation using } m \text{ samples}}$$

## GANs - Training - Considerations

- Optimizers: Adam, Momentum, RMSProp
- Training phase can last for an **arbitrary number** of steps or epochs
- Training is completed when the discriminator is **completely fooled** by the generator.
- The goal of GAN training is to reach a **Nash Equilibrium** where the best D can do is random guessing.

## Type of GANs

Two big families:

- **Unconditional** GANs (just described)
- **Conditional** GANs (2)

## References

---

- [1] Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio, *Generative Adversarial Networks*.
- [2] Mehdi Mirza and Simon Osindero, *Conditional Generative Adversarial Nets*.